
On Choosing a Task Assignment Policy for a Distributed Server System

21.05.2012

ferhat elmas
214805

Outline

- Introduction
 - Motivation
 - Assumptions
 - Performance Metrics
 - System Parameters
 - Simulation
 - Results
 - Confidence Intervals
 - Conclusion
 - Q&A
-

Introduction

- A task assignment policy for a distributed server system that composed of multiple hosts which process tasks in FCFS.
1. **Random:** One of the hosts is randomly chosen. Each of host has an equal probability.
 2. **Round Robin:** Tasks are assigned to hosts in cyclic fashion.
 3. **Dynamic:** The task is assigned to the host that has minimum expected waiting time.
 4. **Size Based:** Task size is partitioned and tasks that are in a certain range are assigned to a particular host.
-

Motivation

- Assignment policies are studied extensively
 - Modeled as exponential, poor model
- Real task values show heavy tail property
 - Very small fraction of task makes nearly half of the load
 - $\Pr\{X > x\} \sim x^{-a}$ where $0 \leq a \leq 2$
 - The lower ***a***, the more variable the distribution

Task Description	a
Unix CPU requirements at BellCore	[1, 1.25]
Unix CPU requirements at Berkeley	$\sim [1]$
Size of files transferred through Web	[1.1, 1.3]
Size of FTP transfers in the internet	[0.9, 1.1]

Assumptions

- Task Size is known in advance
 - Service time is equal to size of task, no scheduling overhead
 - CPU is the only resource
 - Hosts in the server are identical
 - Task is assigned to a host as soon as it arrives.
 - Execution is in FCFS and non-preemptive.
-

System Model

- Task arrival rate is poisson.
 - Task size is bounded pareto distribution.
 - Defined as $B(k, p, a)$
 - $f(x) = a * k^a * x^{(-a-1)} / 1 - (k/p)^a$
 - k and p are smallest and largest task sizes
 - a controls the variance and smaller a means higher variance
 - a varies from 1 to 2
-

Performance Measures

- Average Waiting Time
 - Average SlowDown
 - Throughput
 - Server Utilization
-

System Parameters

- Number of Servers
 - Load (directly affects arrival rate)
 - a , task size variance
 - while a is changing, k is adjusted to keep the mean same
-

Schedule

- Random:
 - Task t is assigned to host i with $1/\text{number_of_hosts}$ probability
 - Round Robin:
 - Task t is assigned to host i ,
 - if $\text{the_number_of_t} \bmod \text{number_of_host} == i$
 - Dynamic:
 - Task t is assigned to host i ,
 - if $\min(\text{host_expected_waiting_time}) == i$
 - Size Based:
 - Task t is assigned to host i ,
 - size_of_t is in $[\text{host_i_lower}, \text{host_i_higher})$
-

Task Ranges

- Ranges are defined such that expected task size within range is equal in each range
 - Each range has a mean of overall mean is divided by number of hosts
 - Cut off values
 - $k = x_0 < x_1 < x_2 < x_3 < \dots < x_{h-1} < x_h = p$
 - if size is in range of $[x_{i-1}, x_i)$, host i is chosen
-

Simulation

- To study variance task size variability
 - a varies from 1.1 to 1.9 with step 0.1
 - p is 100000, mean is 300
 - k is adjusted to keep mean at 300 all time
 - Load is fixed at 0.8
-

Simulation - 2

- To study the number of hosts
 - varies between 1 and 10
 - To study arrival rate
 - utilization(load) is changed
 - $\text{rate} = \text{numberOfServers} * \text{load} / \text{meanTaskSize}$
 - varies between 0.1 and 0.9
-

Simulation - 3

- Number of arrival tasks = 100000
 - Number of runs = 100
 - Steady state after 20000 task arrival
 - important, sufficiently large number of heavy tasks are needed otherwise statistics are opportunistic for policies
-

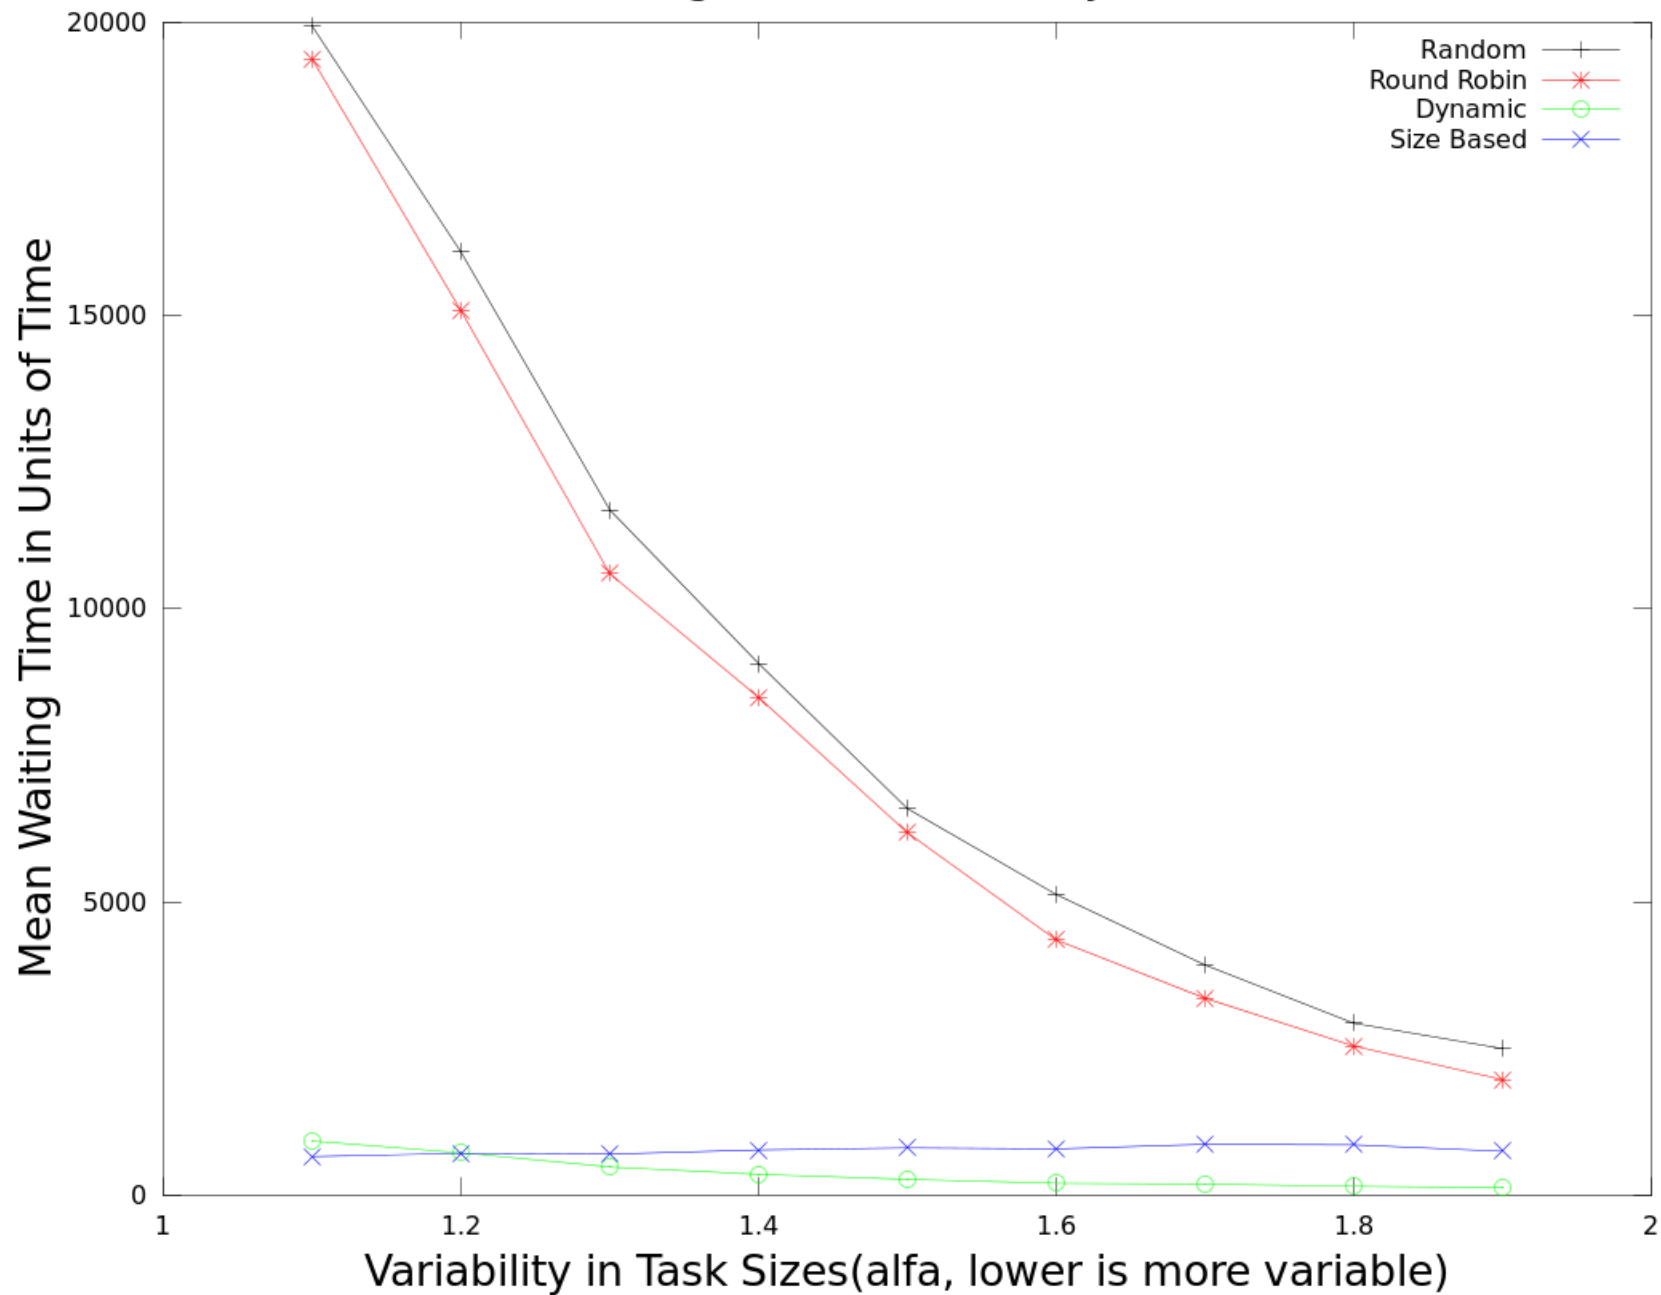
Results

- Variability
 - Number Of Servers
 - Load
-

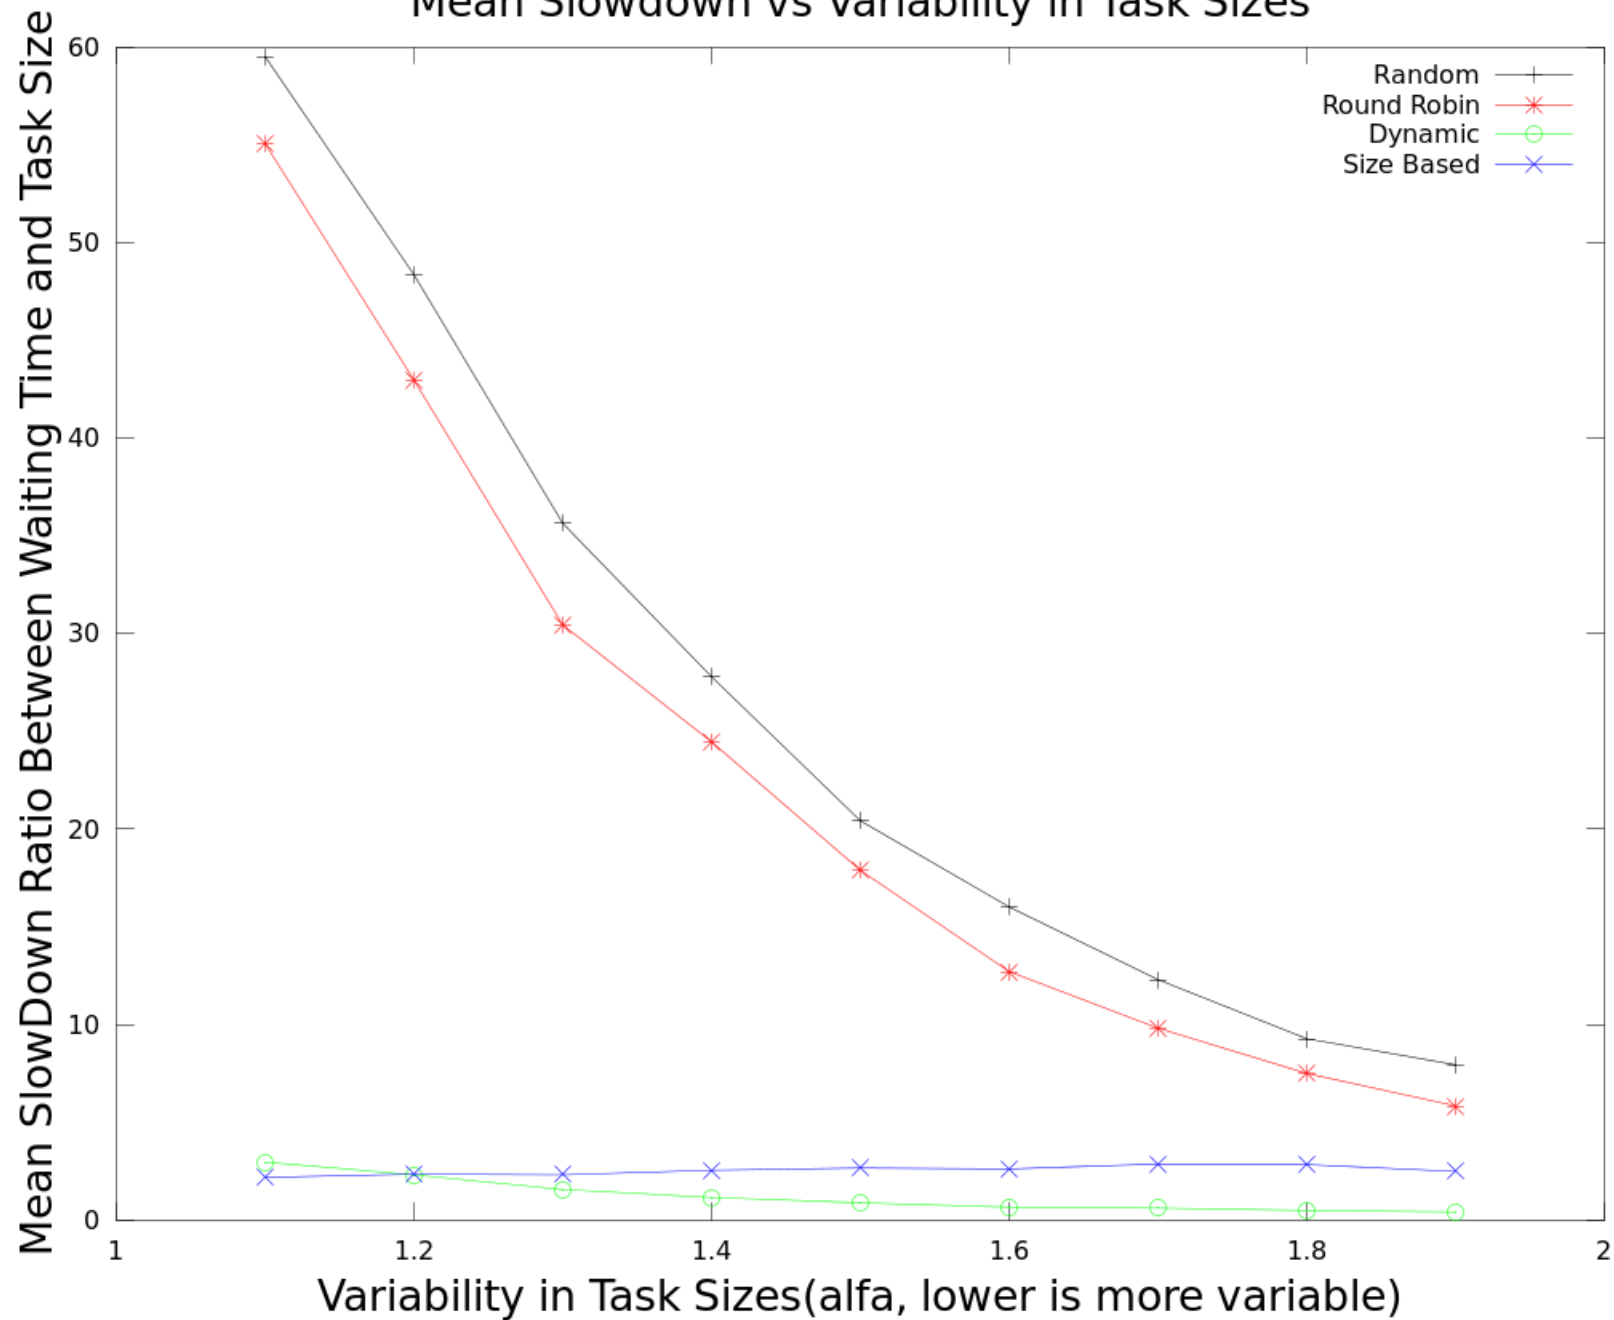
Variability

- a is changed
 - [1.1, 1.9]
- load is fixed at 0.8

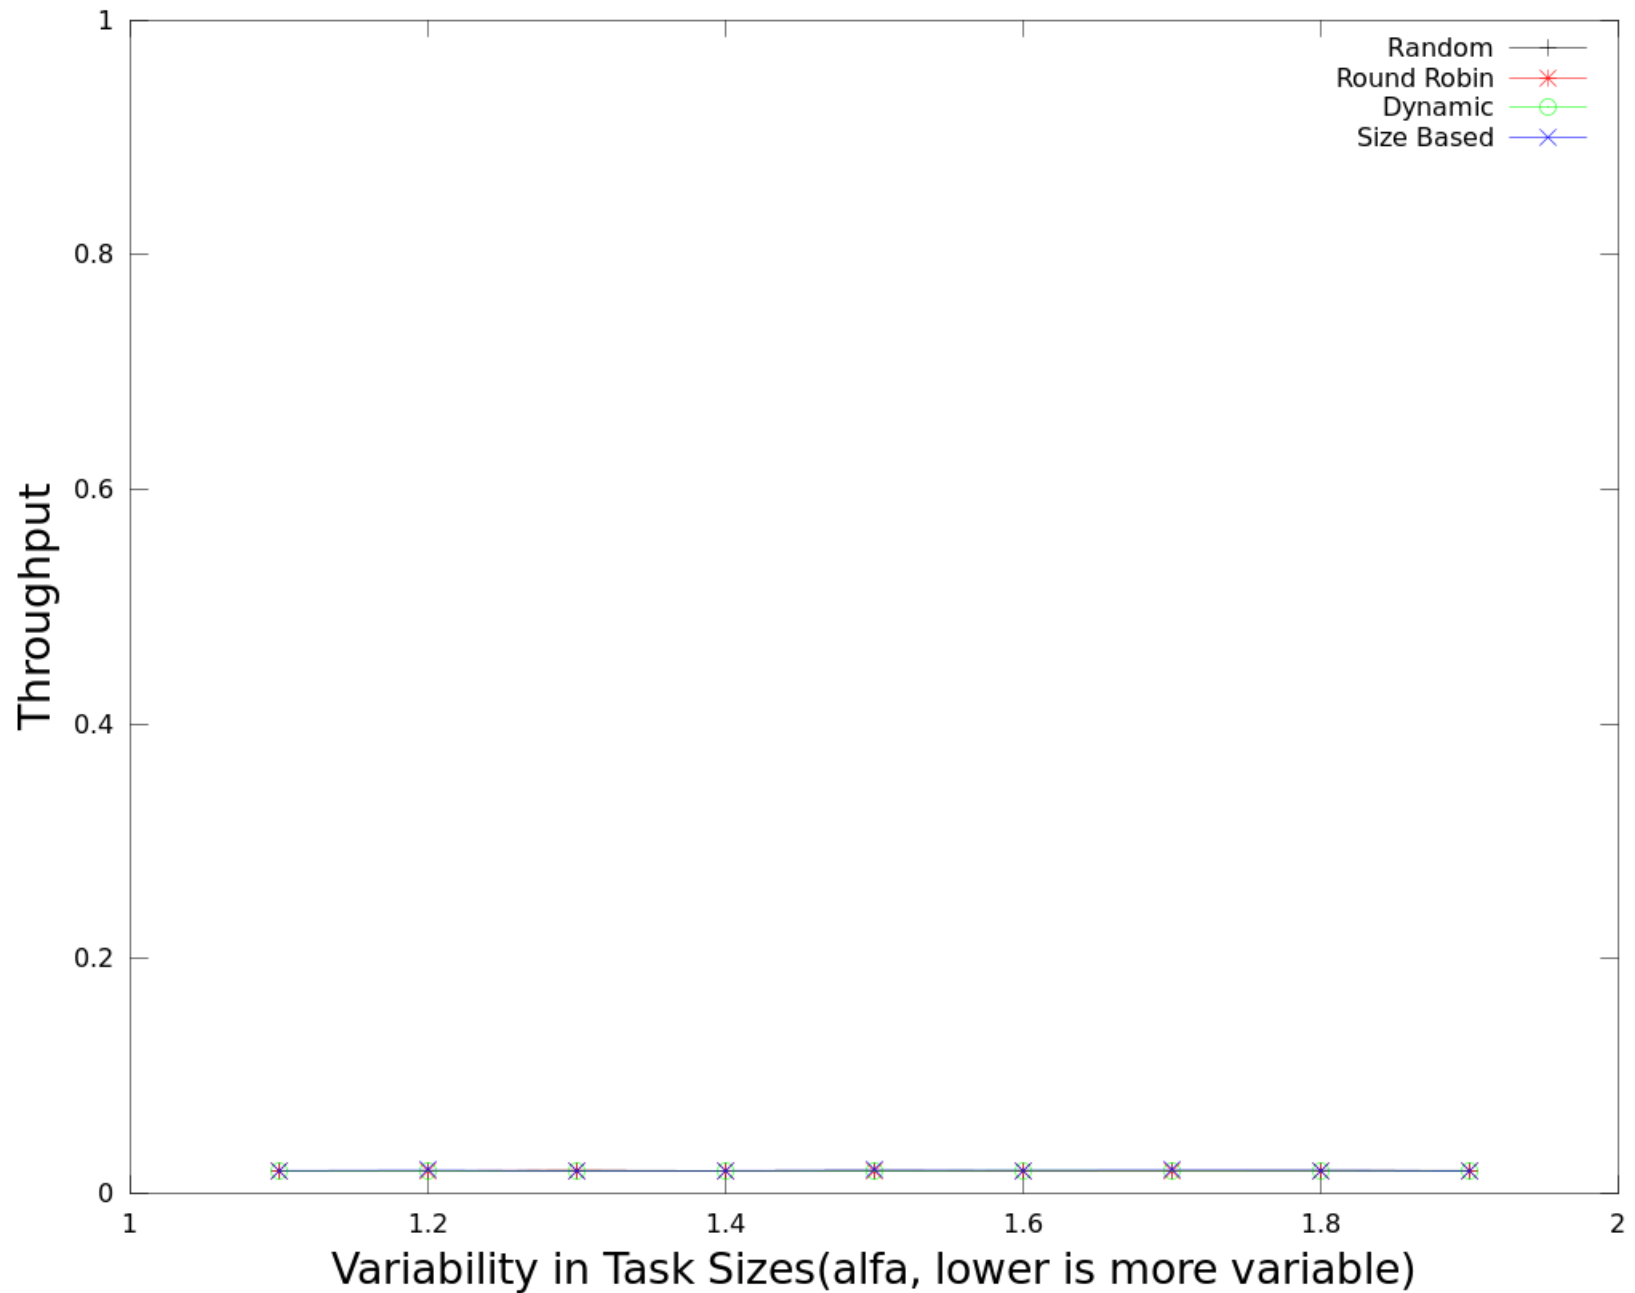
Mean Waiting Time vs Variability in Task Sizes



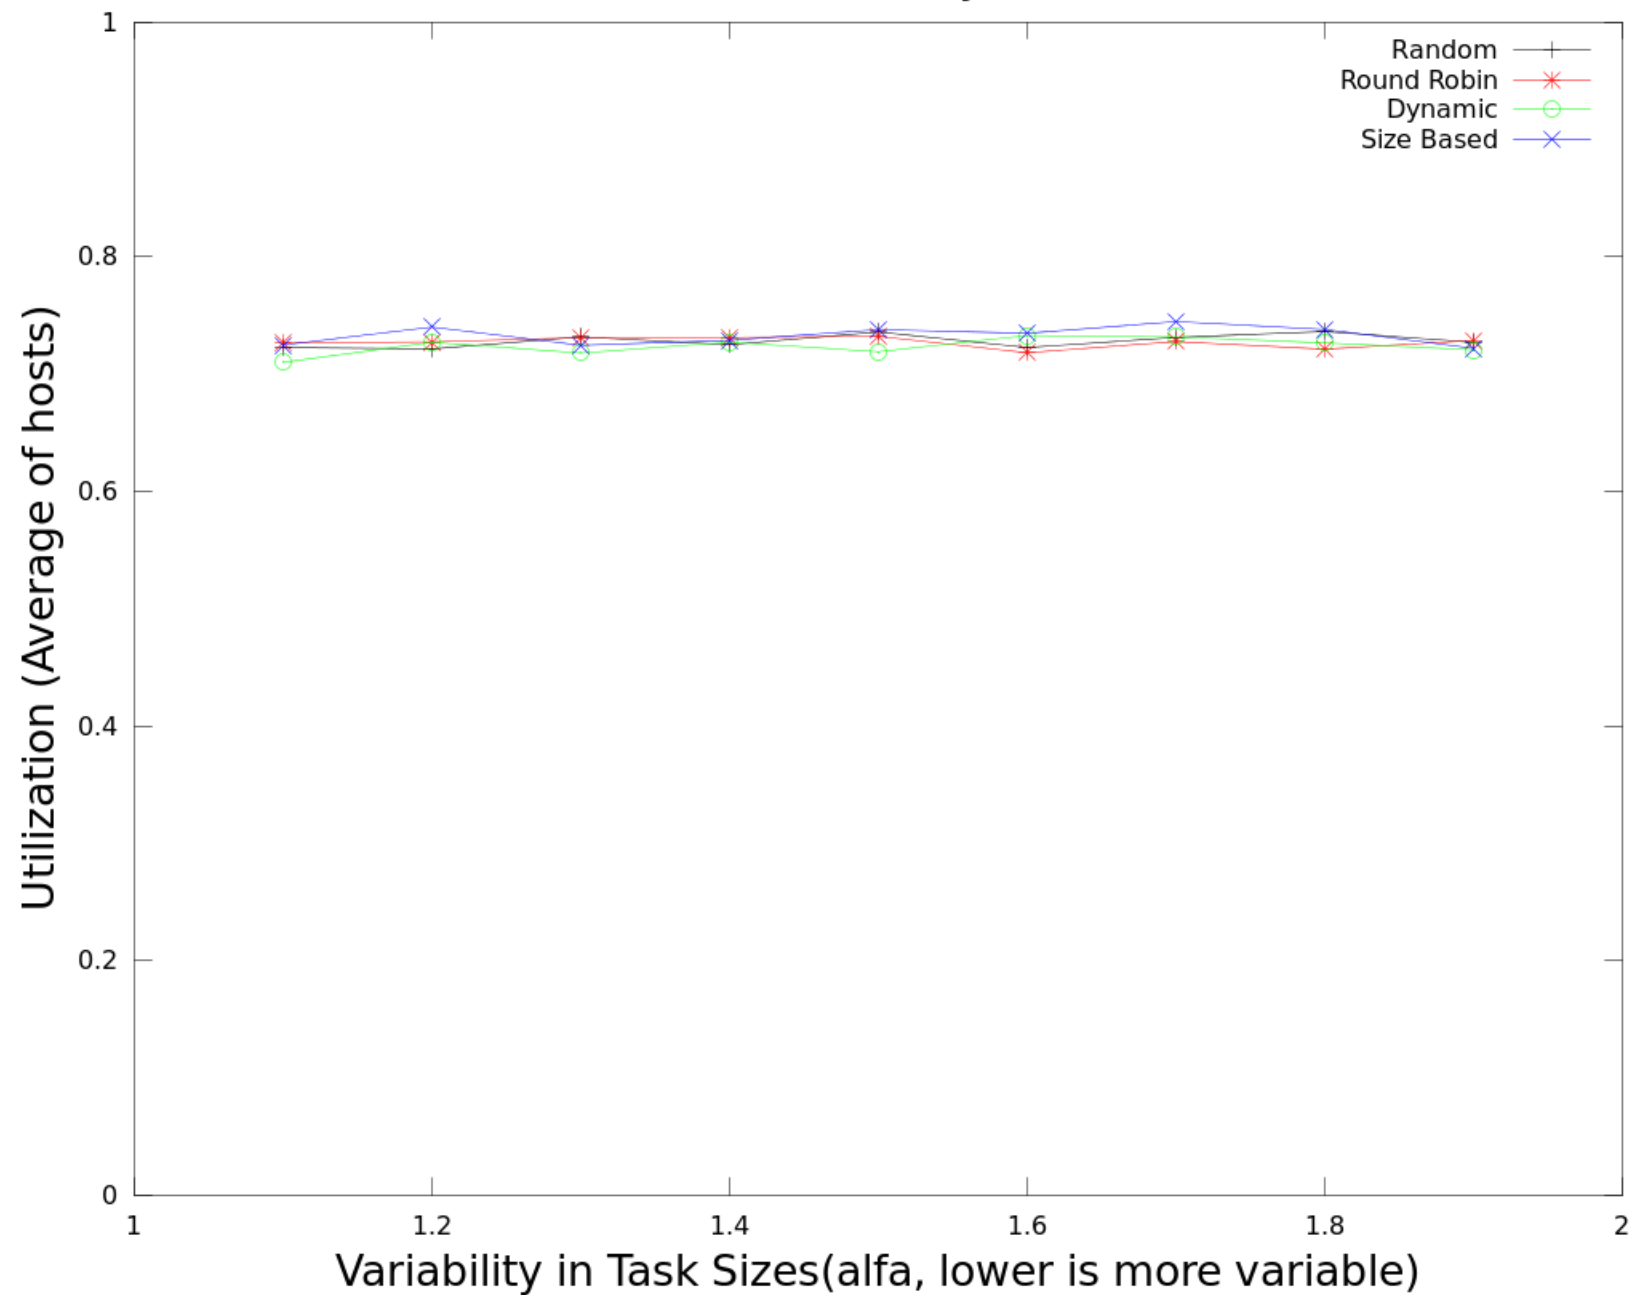
Mean Slowdown vs Variability in Task Sizes



Throughput vs Variability in Task Sizes

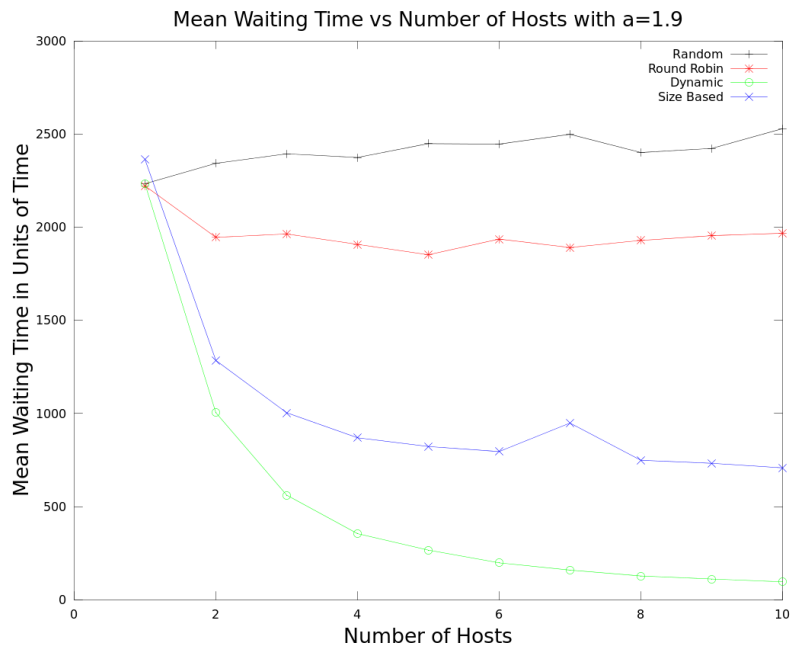
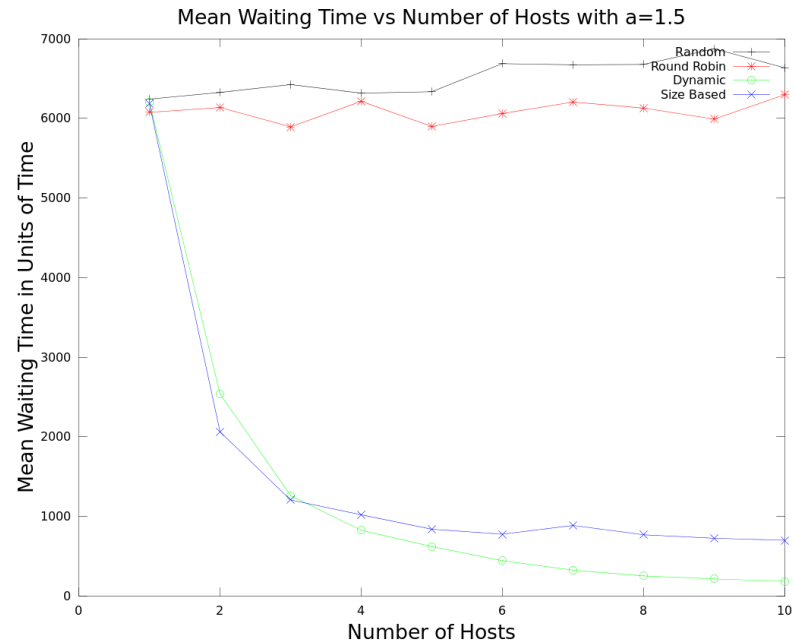



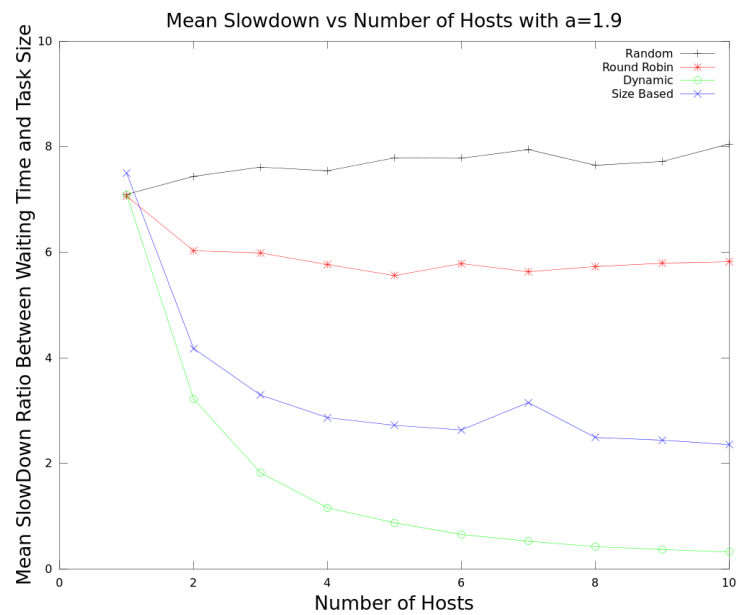
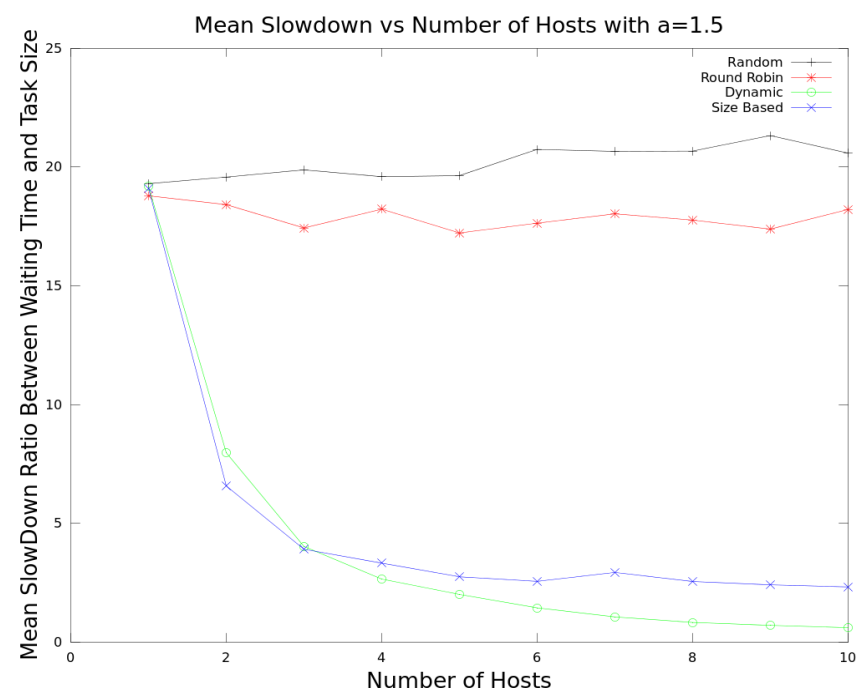

Utilization vs Variability in Task Sizes

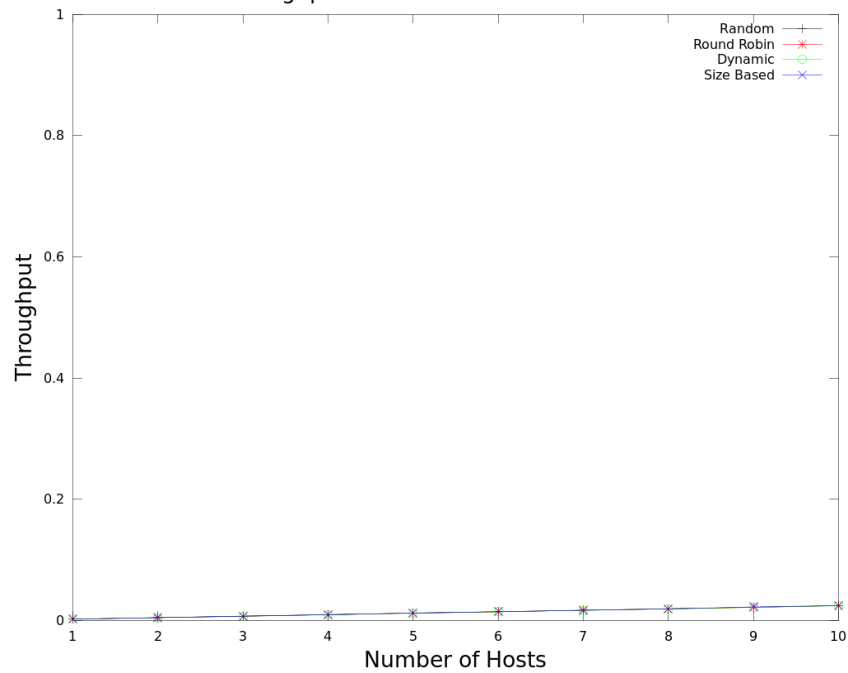





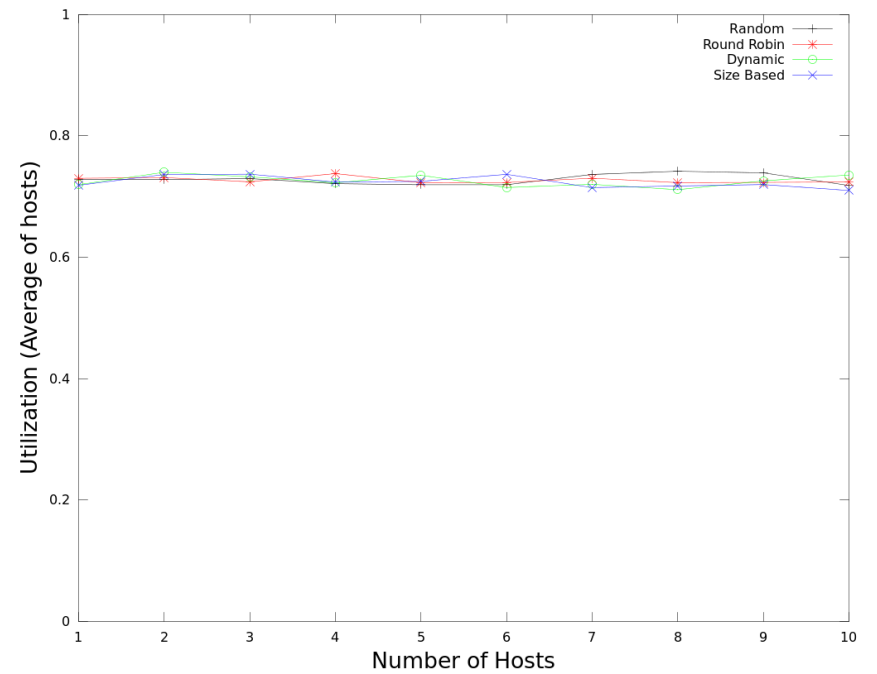

Number Of Hosts

- Number of hosts changes
 - [1, 10]
- Load is fixed at 0.8
- a has 3 different values,
 - 1.1, 1.5, 1.9






Throughput vs Number of Hosts with $a=1.1$ Throughput vs Number of Hosts with $a=1.5$ Throughput vs Number of Hosts with $a=1.9$ 

Utilization vs Number of Hosts with $a=1.1$ Utilization vs Number of Hosts with $a=1.5$ Utilization vs Number of Hosts with $a=1.9$ 

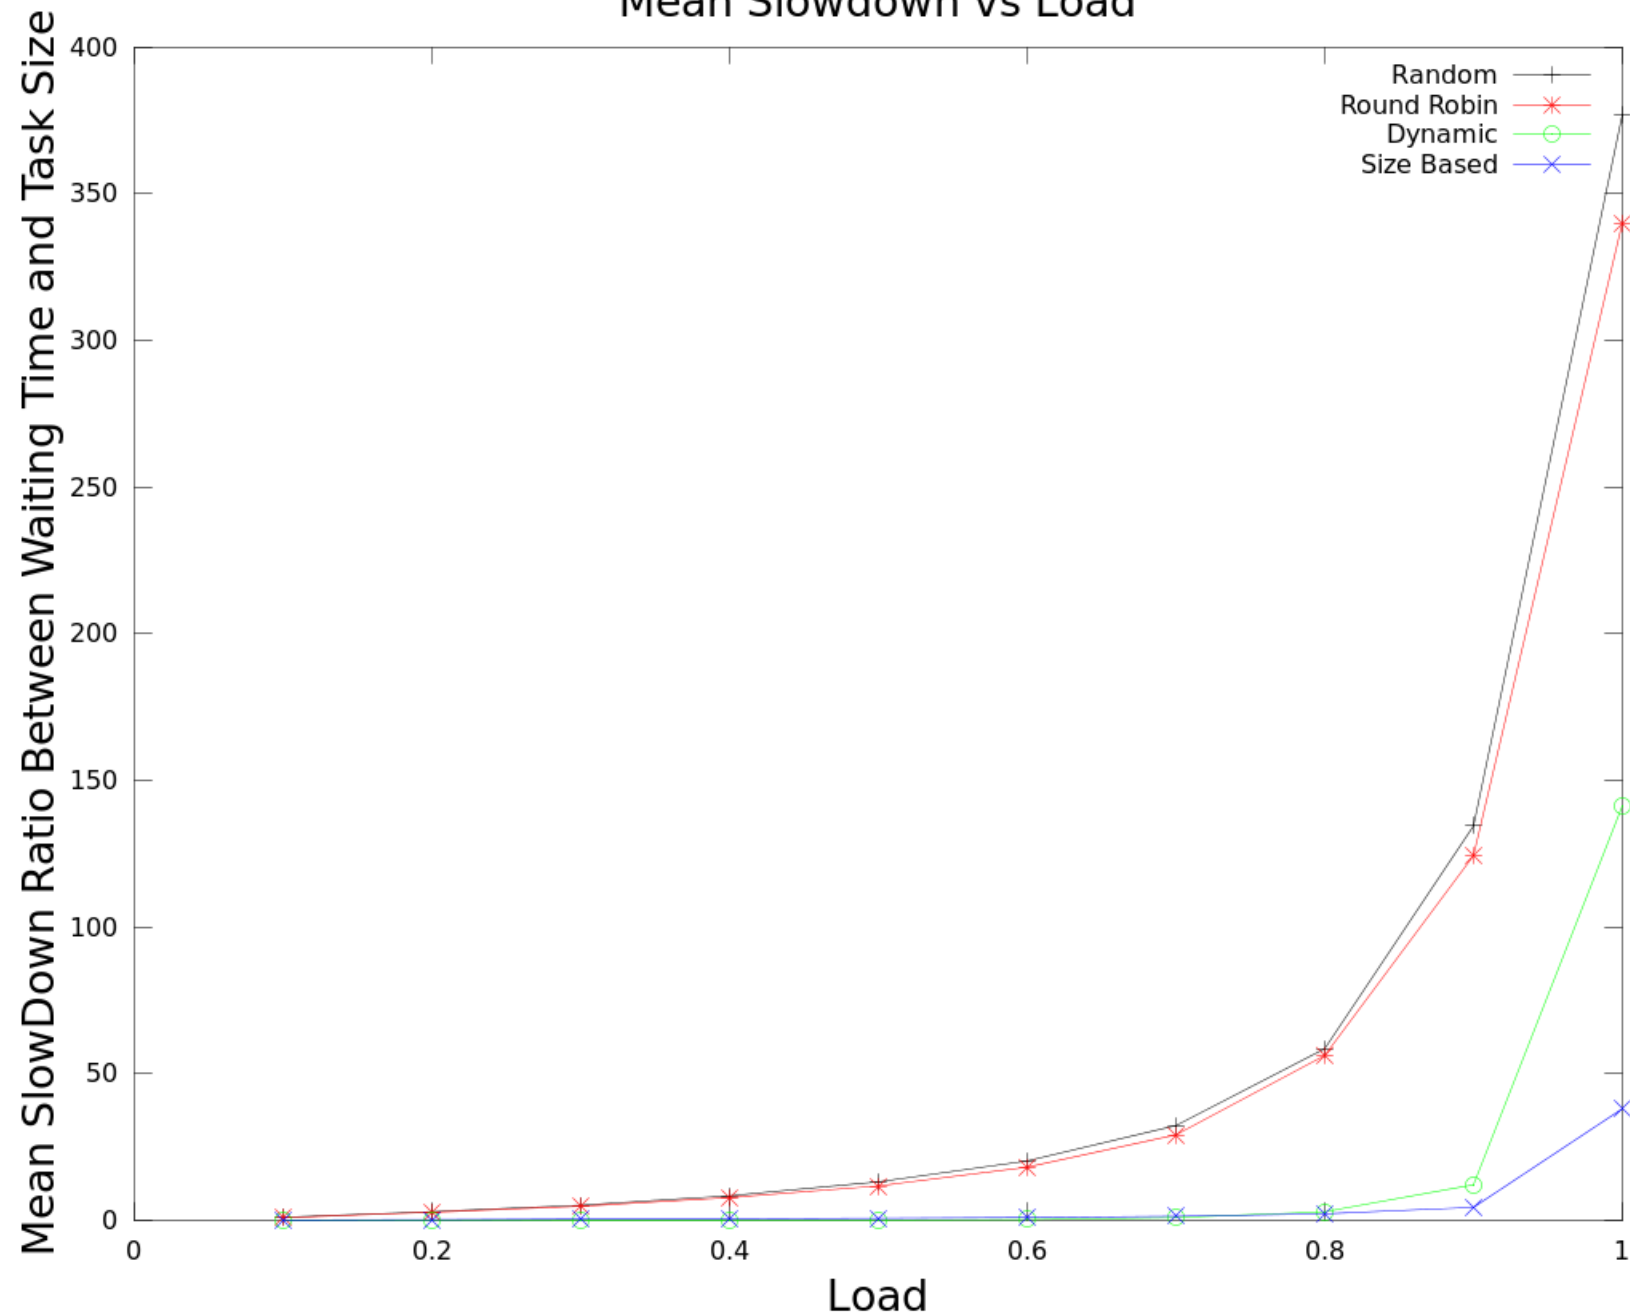
Load

- Varies between 0.1 and 1.0 with step 0.1
 - Variability $a = 1.1$
 - Number of host = 8
-


Mean Waiting Time vs Load




Mean Slowdown vs Load



Throughput vs Load



Utilization vs Load



Conclusion

- With high variance, size based policy does the best
 - With a increases, variability decreases, dynamic policy outperforms others
 - Size-based and dynamic always have much smaller waiting time and slowdown than random and round robin policies
 - Size based policy isn't affected a lot by variability in task size
 - Throughput and utilization are independent of the policy
-

Q & A
